

Research Article

Jon Williamson*

Coaching Aggression, Performance, and Knowledge Transmission in the NFL

<https://doi.org/10.1515/jqas-YYYY-XXXX>

Received November 6, 2024; revised Month DD, YYYY; accepted Month DD, YYYY

Abstract: This study examines coaching aggression and performance in the NFL from 2006-2024. Using machine learning predictions on 643,290 plays, we quantify aggression across four dimensions (fourth down decisions, pass-run balance, deep passing, two-point conversions) and analyze its relationship with Wins Above Replacement. Cross-sectional correlations between aggression and performance declined from $r = 0.197$ ($p = 0.007$) in 2006-2011 to $r = -0.018$ ($p = 0.790$) in 2018-2024. Two-way fixed effects analysis reveals a significant pooled effect: one standard deviation increase in aggression corresponds to 0.235 additional WAR games ($p = 0.020$, Cohen's $d = 0.13$). However, era-stratified analyses are underpowered (13-27% power), preventing conclusions about temporal changes in the causal effect. The cross-sectional erosion reflects diffusion (mean aggression increased significantly, $p < 0.001$) and increased variance (23% rise in standard deviation). Aggression shows moderate year-to-year persistence ($r = 0.458$), but minimal inheritance from mentors to protégés ($r = 0.050$, $p = 0.261$). Fourth down philosophy is the exception, showing modest inheritance ($r = 0.150$, $p < 0.001$) driven entirely by offensive mentors. These findings illuminate how competitive advantages erode through diffusion and have implications for coaching evaluation.

Keywords: NFL coaching, aggression, coaching trees, competitive advantage

1 Introduction

Throughout the 2023 season, Los Angeles Chargers head coach Brandon Staley's aggressive fourth-down philosophy became increasingly controversial (CBS Sports, 2021). In Week 3 against Minnesota, with a precarious four-point lead late in the game, Staley went for it on 4th-and-1 from his own 24-yard line rather than punting (NBC Sports, 2023). The Chargers didn't convert, giving Minnesota excellent field position. Though the Vikings ultimately failed to score, the decision sparked widespread criticism. By December 2023, after a historic 63-21 loss to the Raiders, Staley was fired (Los Angeles Chargers, 2023).

Yet Staley's decisions weren't necessarily reckless. Fourth-down analytics consistently show that aggressive coaches win more games (Romer, 2006; Yam and Lopez, 2019). Teams that "go for it" on fourth down in marginal situations would gain approximately 0.4 additional wins per season on average (Yam and Lopez, 2019). Expected points models demonstrate that punting from midfield often surrenders value (Burke, 2009). So was Staley in the right or wrong? Can coaches be too aggressive?

The answer reveals a fundamental puzzle in competitive strategy: when everyone adopts the same "optimal" approach, does it remain optimal?

NFL coaching represents one of the highest-stakes decision-making environments in professional sports. Head coaches command average salaries exceeding \$6 million annually, with total compensation packages often reaching eight figures (Sportskeeda, 2024). Yet evaluating coaching performance remains remarkably subjective. Teams fire coaches based on win-loss records that confound coaching quality with inherited roster talent, front office decisions, and random variation. Organizations desperately seek objective frameworks for evaluating coaching effectiveness independent of circumstances.

*Corresponding author: Jon Williamson, Knowledge Transmission in the NFL, e-mail: jon@williamsonconsultinggroup.com

This evaluation challenge is particularly acute for aggression—the willingness to deviate from conventional strategy in pursuit of competitive advantage. Should teams go for it on 4th-and-3 from their own 35? Should they pass on 2nd-and-short? Should they attempt two-point conversions with leads? These decisions accumulate into numbers in the win and loss columns each season. But measuring aggression objectively, tracking how it evolves over time, and understanding whether it persists as a stable coaching trait versus adapting situationally has proven elusive. That is the goal of this project.

1.1 Defining Coaching Aggression

We define coaching aggression as the systematic tendency to deviate from conventional strategy in play-calling and game management scenarios. In this project, we believe that aggression can be objectively measured by comparing actual decisions to circumstance-adjusted predictions. We measure aggression across four key dimensions:

- **Fourth-down aggression:** The willingness to attempt conversions rather than punts or field goals. This has become the canonical measure of coaching aggression since Romer’s (2006) influential analysis demonstrated systematic conservatism.
- **Pass-heavy aggression:** The tendency to pass rather than run in neutral situations where defenses expect balanced play-calling. This captures strategic unpredictability and willingness to attack through the air.
- **Deep pass aggression:** The preference for passes traveling beyond the line to gain in passing situations. This represents risk-taking within the passing strategy itself.
- **Two-point aggression:** The willingness to attempt two-point conversions following touchdowns.

Each dimension captures a different aspect of aggressive decision-making. Together, they form a composite aggression score that summarizes a coach’s overall willingness to deviate from conservative strategy, adjusted for the situations that the coach actually faced.

Critically, we measure aggression using machine learning predictions trained on 643,290 plays from 2006–2024. For each decision point, an optimized XGBoost model predicts the baseline “expected” decision given game situation, score, field position, time remaining, and other factors. Aggression is measured as the difference between the actual and predicted percent each coach went with the “aggressive” choice. This approach isolates genuine strategic tendency from forced situational decisions (e.g., trailing teams must be aggressive).

1.2 The Competitive Dynamics Puzzle

On one hand, if aggressive decisions are objectively better than conservative decisions (higher expected points, higher win probability), then the correlation between aggression and performance should be strong. On the other hand, if teams adapt by preparing specifically for aggressive tendencies, or if teams become predictable through over-reliance on aggressive choices, then the advantage may erode or even reverse.

Economic theory suggests competitive advantages from information should erode as knowledge diffuses (Rogers, 2007). Sports analytics research has documented this pattern in baseball, where early adopters of sabermetrics gained advantages that subsequently disappeared as approaches became universal. But the NFL presents a unique context: defenses can adapt to offensive aggression in ways that don’t apply to baseball statistics. When every team goes for it on 4th-and-short, defenses prepare differently. When every team emphasizes passing, defensive strategies evolve to counter. This interactive adaptation may not just erode advantages but potentially reverse them—making extreme aggression predictable and exploitable.

1.3 The Coaching Tree Framework

When NFL teams hire head coaches, pedigree frequently influences decisions. Coordinators from successful organizations often receive premium consideration. The Bill Walsh, Bill Belichick, and Andy Reid coaching lineages receive substantial attention in hiring discussions, with the implicit assumption that working under successful coaches transfers strategic expertise and winning approaches.

This assumption seems intuitive. Coordinators spend years observing great coaches' decision-making processes, learning their strategic philosophies, and presumably absorbing their approaches to game management, aggression, and innovation. If coaching excellence is learnable through apprenticeship, then mentor success should predict protégé success, and strategic tendencies should transmit from mentors to their protégés.

But does the evidence support this narrative? Two testable predictions follow from the coaching tree hypothesis:

Performance inheritance: Protégés of successful coaches should themselves become successful coaches. If great coaches effectively train their coordinators, then mentor performance should positively correlate with protégé performance when those coordinators become head coaches.

Strategic inheritance: Protégés should adopt similar play-calling tendencies to their mentors. If aggressive coaches instill aggressive philosophies in their coordinators, then mentor aggression should correlate with protégé aggression across the four dimensions we measure.

We directly test these predictions in this project.

1.4 WAR as the Measure of Coaching Success

To evaluate whether aggression correlates with coaching performance, and whether successful coaches produce successful protégés, we require a measure of coaching quality that isolates a coach's contribution from the circumstances they inherit. Traditional metrics like win-loss records confound coaching ability with roster talent, strength of schedule, organizational resources, and luck. A coach who goes 10-7 with a mediocre roster may be outperforming a coach who goes 12-5 with elite talent, yet the latter appears more successful by conventional measures.

We employ Wins Above Replacement (WAR), a context-adjusted metric that estimates how many additional wins a coach generates relative to a replacement-level alternative with identical circumstances. Calculated through machine learning predictions that control for roster quality, salary cap allocation, draft capital, injuries, opponent strength, and organizational context, WAR isolates coaching impact from the talent and resources inherited. A coach with +2.0 WAR generated two more wins than expected given their circumstances; a coach with -1.5 WAR cost their team 1.5 wins relative to baseline expectations. This metric parallels established WAR methodologies in baseball and basketball, where player contributions are measured against readily available replacement-level alternatives. Williamson (2026) developed and validated this approach for NFL head coach evaluation using counterfactual machine learning methods, demonstrating strong predictive accuracy ($R^2 = 0.61$ on held-out data) and proper centering around replacement level. The key insight is that WAR provides an objective baseline for comparing coaches across different rosters, eras, and competitive contexts—precisely what's needed to test whether mentor success predicts protégé success independent of the circumstances each inherits.

1.5 Research Questions

This study addresses five interrelated questions about the evolution, stability, transmission, and value of coaching aggression:

1. **How has the relationship between aggression and performance evolved over time?** Specifically, has the correlation strengthened (suggesting optimal strategies winning out), weakened (suggesting

- defensive adaptation), or remained stable? And has the functional form changed from linear (more is better) to non-linear (moderate optimal)?
2. **How stable is aggression within coaches over time?** If aggression represents deeply held strategic beliefs, it should show high year-to-year persistence. If it reflects situational adaptation, persistence should be low. Examining persistence across different aggression components (4th down, pass-heavy, two-point) reveals which aspects reflect traits versus tactics.
 3. **Does mentor performance predict protégé performance?** If coaching excellence is learnable through apprenticeship and coaching trees represent genuine knowledge transmission, then successful coaches should produce successful protégés. We test this directly using context-adjusted performance metrics (WAR) for both mentors and protégés.
 4. **How much does aggression inherit from mentors to protégés?** The coaching tree narrative suggests strategic philosophies transmit through apprenticeship. Testing mentor-protégé correlations in aggression measures directly evaluates whether these transmission effects exist and which dimensions (if any) genuinely pass from mentors to their protégés.
 5. **How do these patterns differ between offensive and defensive coaches?** Offensive coordinators promoted to head coach make offensive decisions central to their expertise. Defensive coordinators must delegate or learn offensive strategy. This natural experiment tests whether aggression reflects learned beliefs (suggesting coordinators of both types should be similar) versus domain expertise (suggesting offensive coaches should show stronger effects).

2 Literature Review

Existing research has documented conservative bias in NFL decision-making and attempted to quantify coaching effectiveness, but several gaps remain. First, no studies have examined how the aggression-performance relationship has evolved temporally as analytics have disseminated and defenses have adapted. Second, aggression persistence within coaches over time has not been systematically analyzed to distinguish traits from tactics. Third, coaching tree effects have not been tested using context-adjusted performance metrics that control for inherited roster quality. Fourth, whether strategic philosophies genuinely transmit from mentors to protégés—and which dimensions show transmission—remains unexamined. Finally, systematic differences between offensive and defensive coaching backgrounds in these relationships have not been explored.

This project addresses these gaps by examining aggression metrics spanning 2006-2024, using context-adjusted WAR to measure both coaching performance and strategic persistence, and testing mentor-protégé relationships while distinguishing offensive from defensive coaching backgrounds.

3 Methods

3.1 Data

Our analysis integrates three complementary datasets spanning nearly two decades of NFL football. The primary dataset consists of play-by-play records for every offensive snap from the 2006 through 2024 seasons, obtained through the `nfl_data_py` package. This yields 643,290 analyzable plays with complete situational information including down, distance, field position, score differential, time remaining, and environmental conditions. The 2006 starting point is determined by data availability, particularly for air yards data required for our deep passing analysis.

The coaching data comes from comprehensive career histories scraped from Pro Football Reference, covering 536 NFL coaches across 104 years of league history. For each coach, we have complete position-by-position career timelines, allowing us to classify coaches by background (offensive vs defensive vs other) and

construct mentor-protégé relationships based on overlapping team assignments. Coach backgrounds are determined by analyzing their coordinator and position coach experience prior to becoming head coaches, using pattern matching on role titles to identify offensive versus defensive specialization. A coordinator is considered a protégé of a head coach if they worked together for at least one season, creating a network of 4,836 documented coaching relationships.

Performance metrics come from a separate Wins Above Replacement (WAR) calculation that estimates each head coach’s contribution to team success relative to a replacement-level alternative (Williamson, 2026). This dataset contains 1,637 coach-year observations, which we merge with our aggression measures on coach name and season. The merge yields 608 coach-year observations where we have both aggression and performance data, covering 123 unique head coaches from 2006-2024. Of these, we successfully classify 605 observations by coach background type: 318 offensive coaches (52.6%), 287 defensive coaches (47.4%).

3.2 Measuring Aggression

We define coaching aggression as systematic deviation from expected decision-making patterns. Rather than using raw frequencies (which confound individual tendencies with situational factors), we train machine learning models to predict baseline behavior and measure how much coaches deviate from these predictions.

For each of four key decision points, we build XGBoost models trained on situational features but explicitly excluding coach identity. This ensures our baseline predictions reflect what the “average” coach would do in identical circumstances, allowing us to isolate individual coaching tendencies. The four components are: (1) fourth down decisions (go vs. punt/field goal), (2) play-calling tendencies (pass vs. run), (3) pass targeting strategy (beyond vs. behind first down marker), and (4) post-touchdown decisions (two-point conversion vs. extra point).

Table 1 summarizes the predictor variables for each model. All models exclude season, week, and coach identity to avoid confounding our analyses of temporal evolution and individual differences.

Tab. 1: Predictive Model Features

Feature Category	Models
Game State (13 features) Score differential, quarter, seconds remaining, timeouts remaining (both teams), down, distance	All
Field Position (5 features) Yard line, goal-to-go, side of field, distance, down	4th Down, Run/Pass, Pass Target (excluded from Two-Point)
Environmental (6 features) Home/away, division game, roof, surface, temp, wind	All
Drive Context (3 features) Drive play count, drive first downs, net yards	All
Vegas Lines (2 features) Spread, total	All
Total Features by Model:	
4th Down Decision	26
Run vs Pass	26
Pass Target Strategy	26
Two-Point Conversion	21

Model performance is strong across all components. The fourth down model achieves an AUC of 0.968, the two-point conversion model reaches 0.927, the run-pass model achieves 0.786, and the pass target model

reaches 0.732. These performance levels indicate that situational factors alone explain the vast majority of decision-making variance, leaving a relatively small but measurable residual that we attribute to coaching philosophy.

For each coach-year, we calculate component-specific aggression scores as:

$$A_c = \text{Actual Rate}_c - \text{Predicted Rate}_c \quad (1)$$

where c indexes the four components. The composite aggression score is the average of the four standardized component scores.

3.3 Performance Analysis

To test whether aggression predicts better coaching performance, we regress average annual WAR on our aggression measures using ordinary least squares. We run separate regressions for composite aggression and each of the four components. To examine temporal evolution, we split the sample into three eras (2006-2011, 2012-2017, 2018-2024) and estimate era-specific relationships. We also calculate rolling 3-year correlations to capture more granular temporal dynamics.

We also analyze offensive and defensive coaches separately to test whether the aggression-performance relationship differs by coaching background. Since all our aggression measures focus on offensive decisions, we hypothesize that offensive coaches should show stronger relationships with performance.

3.4 Era Classification

We classify observations into three approximately equal eras: 2006-2011 (early diffusion, $n=190$), 2012-2017 (mainstream adoption, $n=192$), and 2018-2024 (saturation, $n=224$). This classification is motivated both theoretically (corresponding to distinct phases of analytics adoption in the NFL) and empirically (validated through continuous-year interaction models and Chow tests for structural breaks, detailed in Results).

3.5 Persistence Analysis

To measure the stability of coaching aggression over time, we construct lagged datasets matching each coach-year with their aggression in subsequent years:

$$\rho_k = \text{Corr}(\text{Aggression}_t, \text{Aggression}_{t+k}) \quad (2)$$

for lags $k \in \{1, 2, 3\}$. We estimate these correlations separately for offensive and defensive coaches to test whether coaching background affects the stability of aggressive tendencies. If aggression represents core philosophical beliefs for offensive coaches but situational adaptation for defensive coaches, we should observe stronger persistence among offensive coaches.

3.6 Inheritance Analysis

To test whether aggressive tendencies transmit through coaching relationships, we match coaches with their documented mentors based on career overlap. We analyze coordinator-to-head-coach relationships separately by mentor background (offensive vs defensive) and by coordinator type (OC vs DC).

For each mentor-protégé pair, we calculate average career aggression scores (pooling across all seasons to reduce measurement error) and estimate:

$$\rho_{\text{inherit}} = \text{Corr}(\text{Aggression}_{\text{protégé}}, \text{Aggression}_{\text{mentor}}) \quad (3)$$

Under an inheritance hypothesis with coaching background effects, we would expect substantial positive correlations.

3.7 Multiple Comparison Correction

Our comprehensive analysis involves 98 distinct hypothesis tests across the various dimensions of aggression, performance relationships, persistence, inheritance patterns, within-coach causal effects, and temporal robustness. To address concerns about false discoveries from multiple testing, we apply the Benjamini-Hochberg procedure to control the False Discovery Rate (FDR) at $\alpha = 0.05$ (Benjamini and Hochberg, 1995). This method is more powerful than the conservative Bonferroni correction while still providing strong protection against spurious findings.

The Benjamini-Hochberg procedure ranks all p-values from smallest to largest and identifies the largest rank i where $p_i \leq (i/n) \times \alpha$, where n is the total number of tests. All hypotheses up to this rank are rejected. This approach controls the expected proportion of false discoveries among rejected hypotheses rather than the family-wise error rate.

In our results, we report both raw p-values and indicate which findings survive the Benjamini-Hochberg correction. Of the 63 tests with raw $p < 0.05$, 59 (94%) remain significant after correction, with a cutoff of $p \leq 0.0258$. Our core findings—temporal trends in the aggression-performance relationship (including temporal robustness tests), year-to-year persistence patterns, inheritance from offensive coordinators, and the pooled within-coach causal effect—all survive this correction with highly significant p-values ($p \leq 0.026$). The four tests that lose significance are marginal findings ($p = 0.031$ - 0.043) in subgroup analyses that we interpret cautiously regardless.

3.8 Within-Coach Fixed Effects

To strengthen causal inference, we employ two-way fixed effects regression that controls for both coach-specific quality and temporal changes in aggression’s effectiveness. This approach isolates whether coaches perform better when they are more aggressive than usual, removing confounds from stable coach characteristics and league-wide temporal trends.

The model specification is:

$$\text{WAR}_{it} = \beta \cdot \text{Aggression}_{it} + \alpha_i + \gamma_t + \epsilon_{it} \quad (4)$$

where α_i represents coach fixed effects (controlling for time-invariant coach quality) and γ_t represents year fixed effects (controlling for temporal changes in the value of aggression). We demean variables within both coaches and years, then estimate β using OLS with standard errors clustered by coach.

This two-way fixed effects approach is critical because coaches’ aggression may increase over time as league-wide norms shift. Without year fixed effects, within-coach increases in aggression would be confounded with the declining competitive advantage of aggression, potentially obscuring a genuine causal relationship. The two-way approach asks: when a coach is more aggressive than the league average in year t (relative to their typical deviation from league norms), do they perform better?

4 Results

4.1 Aggression and Performance: A Disappearing Advantage

Our regression analysis reveals that the relationship between coaching aggression and performance is both real and dynamic. We use Wins Above Replacement (WAR) as our performance metric, which controls for

roster quality and inherited team strength by comparing actual performance to expected performance given the roster. Table 2 presents the overall relationships pooling across all seasons from 2006-2024.

Tab. 2: Aggression vs. WAR: Overall (2006-2024)

Measure	r	P-value
Composite	0.081	0.045
Pass-Heavy	0.122	0.003**
4th Down	0.041	0.308
Deep Pass	0.020	0.621
2-Point	-0.021	0.614
N=606 coach-years		
** $p < 0.01$ (survives BH correction)		

Pass-heavy aggression shows a statistically significant weak positive relationship with coaching performance when pooling across all 19 seasons ($r = 0.122$, $p = 0.003$). Composite aggression also shows a positive relationship ($r = 0.081$, $p = 0.045$), though this finding is marginal and does not survive multiple comparison correction. These overall relationships mask temporal evolution. When we split the sample into three approximately equal eras (2006-2011, 2012-2017, 2018-2024), a pattern of decline emerges (Table 3 and Figure 1).

As described in Methods, we applied Benjamini-Hochberg correction to all 98 hypothesis tests in this paper (cutoff: $p \leq 0.0258$). Throughout the Results section, all findings with $p \leq 0.026$ are robust to multiple comparison correction, while marginal results ($p = 0.026$ -0.09) should be interpreted cautiously. The four tests that lose significance after correction are marginal subgroup findings ($p = 0.031$ -0.043) including the overall composite aggression relationship, which we discuss appropriately as suggestive rather than definitive.

Tab. 3: Aggression vs. WAR by Era

Era	Composite (r, p)	Pass-Heavy (r, p)	N
2006-2011	0.197, 0.007**	0.106, 0.147	190
2012-2017	0.095, 0.189	0.138, 0.056	192
2018-2024	-0.018, 0.790	0.102, 0.128	224
** $p < 0.01$			




Fig. 1: Scatterplots of composite aggression (top row) and pass-heavy aggression (bottom row) versus annual WAR, split by era. The strength of the composite aggression relationship has declined from $r = 0.197$ ($p = 0.007$) in the early era to $r = -0.018$ ($p = 0.790$) in the late era.

In the early era (2006-2011), composite aggression shows a statistically significant relationship with performance ($r = 0.197$, $p = 0.007$). Aggressive coaches appeared to outperform their conservative peers. By the middle era (2012-2017), the relationship had weakened ($r = 0.095$, $p = 0.189$). By the modern era (2018-2024), the relationship had weakened to near zero ($r = -0.018$, $p = 0.790$) and is no longer statistically distinguishable from zero. In less than 20 years, aggression went from a modest competitive advantage ($r = 0.20$) to no detectable relationship.

4.1.1 Temporal Robustness

To verify that our era classification is not arbitrary and that the erosion is genuine rather than an artifact of how we divided the time period, we conduct two complementary statistical tests that avoid imposing discrete era boundaries.

First, we estimate a continuous-year interaction model:

$$\text{WAR}_{it} = \beta_0 + \beta_1 \cdot \text{Aggression}_{it} + \beta_2 \cdot \text{Year}_t + \beta_3 \cdot (\text{Aggression}_{it} \times \text{Year}_t) + \epsilon_{it} \quad (5)$$

where Year is centered at 2015 for interpretability. The interaction term β_3 tests whether the aggression-performance relationship changes linearly over time. We find $\beta_3 = -1.865$ ($SE = 0.845$, $p = 0.029$), confirming that the aggression effect erodes significantly and linearly over the study period. The implied aggression effects decline from $\beta = 27.98$ ($p = 0.011$) in 2006 to $\beta = -5.59$ ($p = 0.476$) in 2024—a 33.57-point swing representing complete erosion of the competitive advantage.

Second, we conduct Chow tests for structural breaks at every candidate year from 2009-2019. These tests formally identify the year at which the aggression-performance relationship changes most significantly.

The most statistically significant break occurs at 2011 ($F = 3.63$, $p = 0.027$), with the aggression coefficient declining from $\beta = 27.89$ in the early period (≤ 2011) to $\beta = 1.92$ in the late period (> 2011)—a 93.1% reduction. The second-most significant break occurs at 2012 ($F = 2.93$, $p = 0.054$), with a similar 92.6% decline. These data-driven breakpoints validate our theoretical era boundaries, as the most significant structural break falls precisely at the transition between our early (2006-2011) and middle (2012-2017) eras.

Additional breakpoints at 2013 ($p = 0.103$), 2015 ($p = 0.128$), and 2017 ($p = 0.076$) show similar magnitude declines but weaker statistical significance due to smaller sample sizes in later periods. The clustering of the most significant break around 2011 indicates a genuine regime change rather than gradual drift, and confirms that our findings are not artifacts of arbitrary era classification. These robustness checks establish that the erosion we observe is real and substantial, with the aggression-performance relationship declining by over 90% after 2011 regardless of how we test for temporal change.

4.1.2 Diffusion and Increased Variance

The collapse of the aggression-performance relationship reflects two concurrent dynamics. First, diffusion: mean aggression levels increased as analytics-driven tactics spread through the league. Second, increased variance: rather than converging to a single optimal level, coaches diverged in their approach, with some pushing toward extreme aggression while others remained moderate.

The diffusion claim is statistically robust. Linear regression reveals a significant increasing trend in league-wide average aggression over time ($\beta = +0.125\%$ per year, $R^2 = 0.83$, $p < 0.001$), with time alone explaining 83% of the variance in yearly mean aggression. The increase is substantial: average aggression rose from -0.35% in 2006-2011 to $+1.23\%$ in 2018-2024, a 1.58 percentage point increase ($t = -8.76$, $p < 0.001$, Cohen's $d = 5.37$). One-way ANOVA confirms significant differences across the three eras ($F_{2,16} = 43.75$, $p < 0.001$). Figure 2 visualizes this steady rise in league-wide aggression, with particularly sharp increases after 2015.




Fig. 2: Evolution of coaching aggression over time (2006-2024). The top panel shows mean aggression levels by year with shaded regions indicating composite aggression standard deviation, revealing steady increases. The bottom panel shows each aggression component over time, illustrating increased aggression across most components.

Quantitatively, variance changes differed across aggression components (Table 4). Fourth down and two-point aggression showed substantial variance increases (+28% and +63%), while pass-heavy and deep pass aggression showed slight decreases (-3% and -19%).

Tab. 4: Variance Changes in Aggression Components

Component	Early Era SD (2006-2011)	Late Era SD (2018-2024)	Change
Composite	0.0148	0.0182	+23%
4th Down	0.0213	0.0273	+28%
Pass-Heavy	0.0416	0.0402	-3%
Deep Pass	0.0388	0.0315	-19%
Two-Point	0.0268	0.0437	+63%

This pattern indicates experimentation and divergence rather than convergence. If analytics had identified a single homogeneous optimal aggression level applicable to all teams and contexts, we would expect variance to decrease as coaches converged toward this optimum. Instead, variance increased for composite aggression and two components (4th down and two-point) while decreasing for pass-heavy and

deep pass, suggesting either continued experimentation or that optimal aggression varies by team-specific factors (e.g., roster composition, division strength, quarterback skill).

The mechanism for the relationship decline remains ambiguous. Several hypotheses warrant future investigation:

Defensive adaptation: As aggressive tactics became common, defensive coordinators may have developed counter-strategies, reducing their effectiveness. Extreme predictability may have become exploitable.

Competitive equilibration: The relationship may have simply disappeared because aggression became table stakes rather than differentiator, without extreme aggression becoming actively harmful.

4.1.3 Offensive vs. Defensive Coaches

Since these aggression measures focus on offensive decisions (going for it on 4th down, passing vs running, deep pass targeting, two-point attempts), we examine whether offensive-minded head coaches show different aggression-performance relationships than defensive coaches making offensive decisions.

Tab. 5: Aggression vs. WAR by Coach Type

Measure	Offensive		Defensive	
	r	p	r	p
Composite	0.074	0.192	0.120	0.058
Offensive: $n = 310$, Defensive: $n = 251$				

The relationships between composite aggression and average WAR are weak and mostly non-significant (offensive: $r = 0.074$, $p = 0.192$; defensive: $r = 0.120$, $p = 0.058$). Defensive coaches show a marginally stronger relationship that approaches but does not reach traditional 95% significance. Offensive coaches do not show substantially stronger aggression-performance relationships than defensive coaches.

4.1.4 Within-Coach Fixed Effects: Causal Evidence

The cross-sectional correlations above raise a fundamental question: does aggression cause better performance, or do better coaches (or better teams) simply hire/become more aggressive? To address this, we employ two-way fixed effects regression that controls for both coach quality and temporal trends.

To contextualize effect magnitudes, our sample's WAR distribution (in games per season) has mean $\mu = 0.051$ and standard deviation $\sigma = 1.827$, with interquartile range from -1.054 (25th percentile) to 1.347 (75th percentile). The IQR spans 2.401 games, representing the middle 50% of the performance distribution. The within-coach approach compares each coach to themselves across different years, isolating genuine behavioral effects from selection and omitted variable confounds.

Table 6 presents within-coach fixed effects estimates for the effect of a one standard deviation increase in composite aggression on WAR. The pooled analysis reveals that when coaches increase their aggression by one standard deviation relative to their typical level, they gain 0.235 WAR games ($p = 0.020$, Cohen's $d = 0.13$). This effect persists after controlling for all time-invariant coach characteristics and league-wide temporal changes in aggression's effectiveness. An effect of 0.24 games represents approximately 10% of the IQR in the WAR distribution, corresponding to a percentile shift from the 50th to the 55th percentile—a small but meaningful performance advantage.

Tab. 6: Within-Coach Fixed Effects: Aggression → WAR

Sample	Effect (1 SD)	% IQR	p	Power
Pooled (2006-2024)	0.235**	10%	0.020	65%
Early (2006-2011)	0.258	11%	0.177	27%
Middle (2012-2017)	0.172	7%	0.243	21%
Late (2018-2024)	0.131	5%	0.396	13%

Dependent variable: WAR (games per season)
 Independent variable: Composite aggression (1 SD increase)
 Effect: Gain in WAR games from 1 SD increase in aggression
 % IQR: Effect size as proportion of interquartile range (2.401 games)
 Model: Two-way fixed effects (coach + year), clustered SEs
 Power: Post-hoc statistical power for observed effect
 n = 606 coach-years (123 coaches); Era n: 190 / 192 / 224
 ** $p < 0.05$; Pooled Cohen's $d = 0.13$ (small-to-medium)

The era-stratified results reveal suggestive patterns but suffer from severe power limitations. While point estimates decline from 0.258 games in the early era to 0.131 games in the late era (49% reduction), all era-specific effects are statistically non-significant, and post-hoc power analysis reveals that all three era-specific tests are substantially underpowered (27%, 21%, and 13% respectively). The early-era analysis, despite the largest point estimate, has only 27% power and cannot rule out effects as large as 41 WAR games at conventional power thresholds. The middle and late eras are even more severely underpowered, unable to detect effects smaller than 26 and 24 WAR games respectively at 80% power.

The overlapping confidence intervals and low statistical power mean we cannot definitively conclude the effect differs across eras. The declining point estimates are consistent with erosion but could also reflect sampling variation in underpowered tests. The pooled analysis provides the most reliable estimate, demonstrating a significant causal effect that persists across the full 2006-2024 period with moderate power (65%). While suggestive of erosion, the era-stratified results should be interpreted cautiously given power constraints. Future research with larger samples or longer observation periods would be needed to definitively establish whether the causal effect has diminished over time or remains stable.

4.2 Aggression as Trait and State

Understanding whether coaching aggression represents stable individual traits or adaptive responses to circumstances has important implications for talent evaluation. Our persistence analysis examines year-to-year correlations in coaching behavior.

Table 7 presents the year-over-year (lag 1) correlations for each aggression component, based on 467 consecutive coach-year pairs where the same coach was in place for two consecutive seasons. Note that these estimates may be upward-biased due to survivorship: coaches who retain their positions across multiple seasons are disproportionately successful, and successful coaches may have more stable situations that enable consistent strategic approaches.

Tab. 7: Year-to-Year Aggression Persistence (Lag 1)

Component	r	R^2
Composite	0.458***	0.210
4th Down	0.556***	0.309
Pass-Heavy	0.473***	0.224
Deep Pass	0.455***	0.207
2-Point	0.393***	0.154

*** $p < 0.001$

All five measures show statistically significant and practically meaningful persistence. The correlations range from 0.39 to 0.56, indicating that coaching aggression in year N predicts 15-31% of the variance in year $N+1$ aggression. This is substantial enough to constitute a real trait, yet modest enough to leave 70-85% of the variance driven by contextual factors.

Fourth down decision-making shows the strongest persistence ($r = 0.556$, $R^2=0.309$). A coach who goes for it frequently on fourth down one year is quite likely to do so again the next year, even as rosters turn over and circumstances change. This suggests fourth down philosophy represents genuine strategic preferences—core beliefs about optimal decision-making that persist across contexts.

Figure 3 shows how these correlations decay as we extend the lag. At lag 2 (two years apart), correlations drop to the 0.23-0.44 range. At lag 3, they fall further to 0.28-0.40. Remarkably, even at three years apart, all correlations remain statistically significant.




Fig. 3: Persistence of coaching aggression over 1, 2, and 3 year lags. Fourth down aggression (purple squares) shows the strongest and most stable persistence, while two-point aggression (green triangles) decays most rapidly. All measures remain statistically significant even at 3-year lags.

4.2.1 Offensive vs. Defensive Coaches

Persistence patterns differ dramatically by coaching background. Table 8 shows correlations at 1-, 2-, and 3-year lags separated by coach type.

Tab. 8: Persistence by Coach Background (Lags 1-3)

Component	Offensive			Defensive		
	Lag 1	Lag 2	Lag 3	Lag 1	Lag 2	Lag 3
Composite	0.521***	0.460***	0.355***	0.424***	0.387***	0.361***
4th Down	0.690***	0.619***	0.549***	0.373***	0.236**	0.269**
Pass-Heavy	0.544***	0.422***	0.425***	0.448***	0.254**	0.234*
Two-Point	0.424***	0.417***	0.488***	0.425***	0.189*	0.183
Offensive: $n_{lag1} = 228$, $n_{lag2} = 160$, $n_{lag3} = 110$; Defensive: $n_{lag1} = 194$, $n_{lag2} = 150$, $n_{lag3} = 113$						
* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$						

Offensive coaches show much stronger year-to-year persistence than defensive coaches, particularly for 4th down decisions ($r = 0.690$ vs $r = 0.373$). The gap is substantial: offensive coaches explain 48% of next year's 4th down aggression from this year's value, while defensive coaches explain only 14%. For composite aggression, the gap is more modest ($r = 0.521$ vs $r = 0.424$, difference=0.097) but still substantive.

This pattern extends across longer lags. At 2-year lags, the 4th down persistence gap reaches its maximum ($r = 0.619$ vs $r = 0.236$, difference=+0.384). At 3-year lags, large gaps persist for 4th down and pass-heavy aggression, though two-point aggression becomes non-significant for defensive coaches.

The interpretation is clear: for offensive coaches, aggression—especially 4th down philosophy—is a deeply ingrained trait reflecting core beliefs about optimal strategy. These beliefs persist strongly across contexts. For defensive coaches making offensive decisions, aggression is more contextual and reactive, adapting more to personnel, opponents, and situations rather than expressing stable philosophical commitments.

4.3 Testing the Theory of Coaching Trees

The concept of coaching “trees”—lineages descending from influential mentors and carrying forward their strategic philosophies—is deeply embedded in NFL culture. Bill Belichick's disciples, the Andy Reid coaching tree, the Bill Walsh West Coast offense, these narratives dominate hiring discussions and suggest meaningful transmission of strategic DNA.

Our data allow us to test whether from the lens of context-adjusted performance and aggression, coaching trees exist. We analyze 517 documented mentor-protégé pairs where we have aggression data for both mentor and protégé, using career average aggression scores for each coach to maximize reliability and sample size. Table 9 presents results for the full sample.

Tab. 9: Aggression Inheritance: Overall Mentor-Protégé Pairs

Component	<i>r</i>	P-value
Composite	0.050	0.261
4th Down	0.150	0.0006***
Pass-Heavy	0.023	0.596
Deep Pass	0.072	0.101
Two-Point	0.085	0.053*
N=517 mentor→protégé pairs		
* $p < 0.1$		
*** $p < 0.001$		

Composite aggression shows essentially zero inheritance ($r = 0.050$, $p = 0.261$). Pass-heavy aggression is nearly zero ($r = 0.023$, $p = 0.596$). Deep pass targeting shows a weak non-significant positive correlation ($r = 0.072$, $p = 0.101$). Knowing how aggressive a coach's mentor was tells you almost nothing about how aggressive that coach will be. The coaching tree hypothesis, for overall strategic aggression, finds no support.

Fourth down decision-making shows modest but significant inheritance ($r = 0.150$, $p < 0.001$). While explaining only 2.3% of variance, it's statistically robust and aligns with the persistence findings showing 4th down philosophy is the most stable trait.

4.3.1 Offensive vs. Defensive Mentors

The inheritance story becomes much richer when we separate by mentor background. Table 10 shows correlations split by whether the mentor had offensive or defensive coaching background.

Tab. 10: Inheritance by Mentor Background

Component	Offensive		Defensive	
	<i>r</i>	<i>p</i>	<i>r</i>	<i>p</i>
Composite	0.068	0.274	0.042	0.523
4th Down	0.217***	0.0005	-0.016	0.808
Pass-Heavy	0.031	0.618	0.035	0.597
Deep Pass	-0.029	0.649	0.075	0.250
Two-Point	0.099	0.114	0.090	0.170
Offensive: $n = 257$, Defensive: $n = 236$				
*** $p < 0.001$				

Offensive mentors show significant 4th down inheritance ($r = 0.217$, $p = 0.0005$)—notably stronger than the overall sample ($r = 0.150$). Defensive mentors show essentially zero ($r = -0.016$, $p = 0.808$). The difference is substantial in correlation magnitude. For pass-heavy aggression and composite aggression, neither group shows significant inheritance, though offensive mentors show weak positive correlations while defensive mentors show essentially zero correlations.

We can push this further by classifying mentors not just by overall background but by whether they were offensive coordinators (OC) or defensive coordinators (DC). Table 11 shows these results.

Tab. 11: Inheritance by Coordinator Type

Component	OC → HC		DC → HC	
	r	p	r	p
Composite	0.079	0.276	0.041	0.526
4th Down	0.191**	0.008	-0.016	0.808
Pass-Heavy	0.200**	0.005	0.040	0.533
Deep Pass	0.062	0.390	0.064	0.325
Two-Point	0.101	0.162	0.084	0.193

OC: $n = 193$, DC: $n = 241$
 ** $p < 0.01$

Offensive coordinators who become head coaches pass on both pass-heavy tendencies ($r = 0.200$, $p = 0.005$) and 4th down philosophy ($r = 0.191$, $p = 0.008$). The correlations are similar in magnitude, both significant at the 1% level. These correlations explain 4.0% and 3.6% of variance respectively—modest but statistically robust effects.

Defensive coordinators who become head coaches show no statistically detectable inheritance. All correlations are near zero and nonsignificant (4th down: $r = -0.016$, $p = 0.808$; pass-heavy: $r = 0.040$, $p = 0.533$). They don't transmit offensive philosophies because offensive decisions were never core to their coaching identity. Working under a successful defensive head coach teaches adaptation and leadership skills, but not offensive tactical philosophy.

Figure 4 visualizes these patterns with scatter plots split by mentor type.




Fig. 4: Mentor-protégé correlations split by mentor background. Only offensive mentors (orange) show clear positive slopes for 4th down and pass-heavy aggression. Defensive mentors (blue) show flat relationships across all components.

4.3.2 The Null Findings: Deep Pass and Two-Point Aggression

Two aggression components—deep pass targeting and two-point conversion attempts—show essentially null relationships across all analyses. Neither correlates meaningfully with performance (deep pass: $r = 0.013$, $p = 0.754$; two-point: $r = -0.006$, $p = 0.887$), neither shows meaningful inheritance even from offensive coordinators (deep pass: $r = -0.001$, $p = 0.994$; two-point: $r = 0.119$, $p = 0.091$), and while both show year-to-year persistence comparable to other components, the lack of performance relationships and inheritance patterns suggests these may not represent stable coaching philosophies in the same way fourth down and pass-heavy decisions do.

Several explanations warrant consideration. First, these decisions may be too rare to measure reliably—two-point attempts occur infrequently, and deep pass targeting classifications may be noisy. Second, these decisions may be more situation-dependent than coach-dependent: deep pass targeting likely depends heavily on quarterback arm strength and receiver personnel, while two-point decisions in close games may follow league-wide optimal strategies with little room for philosophical variation. Third, these measures may have poor construct validity for capturing meaningful strategic variation. The null findings suggest that not all aspects of offensive decision-making represent transmissible coaching philosophies or stable individual traits.

5 Discussion

5.1 The Erosion of a Competitive Advantage

Our findings document the substantial erosion of a competitive advantage in professional sports, from measurable benefit to no detectable relationship in less than two decades. This pattern mirrors technology adoption in other competitive markets—financial markets arbitraging away trading strategies, military tactics becoming obsolete as adversaries adapt, business strategies eroding through imitation.

In the early period (2006-2011), aggressive tactics provided measurable performance benefits ($r = 0.197$, $p = 0.007$ for composite aggression). But by 2018-2024, the relationship was no longer statistically distinguishable from zero ($r = -0.018$, $p = 0.790$).

The decline reflects two concurrent dynamics. First, diffusion: as analytics-driven decision-making spread through the league, more coaches adopted aggressive tactics. This claim is supported by the statistically significant increase in average aggression between the periods. Second, increased variance: rather than converging to a single optimal level, variance in aggression increased 23% from early to late era. Some coaches pushed toward extreme aggression while others adopted moderate approaches, creating more diversity rather than homogeneity.

- **Defensive adaptation:** As aggressive tactics became common, defensive coordinators may have developed counter-strategies, reducing effectiveness. Predictability may have become exploitable.
- **Competitive equilibration:** Aggression may have simply become table stakes—necessary but no longer differentiating.

The strategic frontier has likely moved to areas we have not yet successfully measured, perhaps scheme innovation in play design, within-game adaptation speed, or player development efficiency.

5.2 Causality: Evidence from Within-Coach Variation

While our cross-sectional correlations demonstrate that aggressive coaches performed better in the early era, establishing causality requires addressing confounding from coach quality and selection effects. Our two-way fixed effects analysis provides evidence that the relationship is at least partly causal.

The within-coach fixed effects approach controls for all time-invariant coach characteristics (ability, personality, leadership style) and temporal trends (league-wide changes in aggression's effectiveness). By comparing each coach to themselves across different years, we isolate whether coaches perform better when they are more aggressive than usual. The significant pooled effect shows that a one standard deviation increase in aggression corresponds to 0.235 additional WAR games ($p = 0.020$, Cohen's $d = 0.13$), even after removing all between-coach differences and temporal trends. This effect represents approximately 10% of the IQR in the WAR distribution, a small but meaningful performance advantage.

This causal evidence strengthens our interpretation of the cross-sectional patterns. The relationship between aggression and performance is not merely that better coaches happen to be more aggressive—within the same coach, being more aggressive in a given year predicts better performance in that year. The pooled effect demonstrates a persistent causal relationship across the full 2006-2024 period.

However, two critical limitations constrain our ability to characterize temporal patterns in this causal effect. First, era-stratified analyses suffer from severe power limitations. While point estimates decline from 0.258 games in early to 0.131 games in late eras (49% reduction), post-hoc power analysis reveals dramatically insufficient power: 27%, 21%, and 13% respectively. With such low power, these tests cannot distinguish true erosion from sampling variation. All era-specific effects are statistically non-significant, and we cannot rule out effects as large as 25-43 WAR games in any single era. Claims of erosion based on within-coach variation must therefore be tempered by these power constraints. Second, within-coach variation in aggression is smaller than between-coach variation, limiting effect detection. Third, reverse causation within seasons remains possible, though situational controls partially address this.

The convergent evidence from multiple approaches—cross-sectional correlations strongest when tactics were rare, erosion as tactics spread, and within-coach effects controlling for stable characteristics—collectively supports a causal interpretation while acknowledging remaining ambiguities. For practitioners, the key insight holds regardless: whether through causal mechanisms or spurious correlation, aggressive coaching no longer associates with better outcomes in the modern era.

5.3 Why Fourth Down Philosophy Transfers

Fourth down decision-making is unique among the aggression components we measure. It's both the most persistent trait over time ($r = 0.56$ year-to-year) AND the most heritable trait across mentor-protégé relationships ($r = 0.19$ - 0.23 for offensive mentors). This dual finding suggests it represents something fundamental about coaching philosophy that other tactical decisions do not.

Why might this be? Several mechanisms seem plausible and warrant deeper investigation:

- **Clarity of analytics.** The research on optimal fourth down decision-making is clearer and more settled than analytics on pass-run balance or deep pass targeting. Expected points models provide relatively unambiguous recommendations for many fourth down scenarios. This clarity may make fourth down philosophy more teachable and transferable as an abstract principle rather than context-dependent judgment.
- **Roster independence.** Unlike passing frequency (which may require a capable quarterback) or deep pass targeting (which may require strong receivers), fourth down philosophy is less constrained by roster construction. Any team can choose to go for it on fourth down; execution may vary, but the decision itself is relatively roster-independent. This independence may allow fourth down beliefs to transfer across contexts more readily.

These mechanisms have significant practical implications. If fourth down philosophy is the one truly transferable element of coaching DNA, organizations should probe this specific dimension carefully in interviews and evaluations. Ask candidates not just what fourth down decisions they would make, but why. The reasoning process may reveal more about transferable philosophy than the decisions themselves.

Conversely, organizations should be skeptical of assumptions that other aspects of a candidate's previous success will transfer. Pass-run balance is heavily context-dependent and unlikely to replicate in

new environments. The coaching tree narrative is mostly myth, with fourth down philosophy as the lone exception that proves the rule.

5.4 Challenging the Coaching Tree Narrative

The coaching tree narrative is pervasive in NFL culture and shapes major hiring decisions. When teams hire a candidate from a successful coaching lineage, the implicit assumption is that successful strategic DNA will transfer. Our data show this assumption is largely false.

The reality is that coaching aggression philosophies are primarily adaptive and context-dependent rather than inherited traits. The correlation between mentor and protégé aggression is essentially zero for composite aggression ($r = 0.06$), pass-heavy aggression ($r = 0.01$), deep pass aggression ($r = 0.04$), and two-point aggression ($r = 0.08$).

Why does the myth persist despite the evidence? Several cognitive biases seem likely:

- **Availability bias.** Memorable successes from coaching trees (Bill Walsh → Mike Holmgren → Andy Reid) become readily available examples that dominate thinking, while numerous failures (which outnumber successes) are forgotten.
- **Narrative convenience.** The coaching tree story provides a compelling narrative structure for explaining success and failure. It's more satisfying to attribute outcomes to lineage than to context, luck, and complex interactions.
- **Confirmation bias.** When protégés of successful coaches succeed, we credit the coaching tree. When they fail, we explain it away as poor fit or circumstances, preserving the core belief.
- **Selection effects.** Successful coaches' protégés get more opportunities to become head coaches, creating a spurious association between pedigree and success that reflects opportunity rather than transferred capability.

Organizations should be explicitly skeptical of pedigree-based hiring.

5.5 Where Has the Advantage Moved?

If aggressive play-calling no longer provides competitive advantages, where has the frontier moved? Organizations seeking edges need specific hypotheses about tomorrow's inefficiencies, not optimization of yesterday's strategies.

Several candidates warrant investigation:

- **Within-game adaptation speed.** If everyone starts games with similar optimal strategies, advantages may accrue to teams that adjust fastest to unfolding game dynamics. This requires real-time analysis capabilities, rapid communication systems, and coaches who can process information and update beliefs quickly rather than rigidly executing predetermined plans.
- **Scheme innovation beyond decision frequency.** Our analysis measures WHEN teams do things (fourth down, two-point, pass vs run) but not HOW. Two teams might pass with identical frequency but run completely different passing concepts. Innovation in play design, route combinations, blocking schemes, and personnel groupings could provide advantages even as decision frequencies converge.
- **Player development efficiency.** Teams that turn draft picks into productive players faster gain cumulative advantages over multiple seasons. This requires coaching staffs skilled not just in tactical decision-making but in teaching, communication, motivation, and individualized skill development.

The broader lesson is this: competitive advantage in professional sports is temporary by nature. What works today will be copied tomorrow. Success requires continually identifying new inefficiencies before they become obvious. In the NFL's arms race of strategic innovation, the critical skill is not aggressive execution of known strategies—everyone does that now—but correctly identifying what will be optimal next.

5.6 Limitations

Several limitations warrant discussion beyond the causality concerns addressed earlier.

First, our play-by-play data begin in 2006, missing earlier eras. We cannot determine whether the early aggressive advantage we observe in 2006-2011 was itself a continuation of an even longer trend or represented a genuine inflection point. Historical data would reveal whether similar adoption curves occurred for previous innovations.

Second, our inheritance analysis is limited by sample size. With 509 coordinator-to-head-coach transitions, we may lack statistical power to detect small inheritance effects. However, the substantively tiny correlations ($r < 0.10$ for most components) suggest that even if such effects exist, they're small enough to be practically unimportant for hiring decisions.

Third, we cannot observe coaching constraints. Some coaches may want to be aggressive but face organizational pressure to be conservative. Others may have coordinator partnerships that override their personal preferences. Our measures capture observed behavior rather than underlying intentions, which may differ due to unmeasured constraints.

6 Conclusion

Through comprehensive analysis of 643,290 plays and 606 coach-years spanning nearly two decades (2006-2024), we document the substantial erosion of aggressive coaching as a competitive advantage in professional football. Our findings challenge widely-held beliefs about coaching decision-making, strategic innovation, and talent evaluation in the NFL while providing a rare empirical window into how competitive advantages disappear through diffusion and adaptation.

Aggressive coaching caused measurable performance benefits in the early analytics era, and two-way fixed effects analysis establishes a persistent causal relationship. Controlling for coach quality and temporal trends, a one standard deviation increase in aggression corresponds to 0.235 additional WAR games ($p = 0.020$, Cohen's $d = 0.13$)—representing approximately 10% of the interquartile range in the WAR distribution, a small but meaningful performance advantage. This pooled effect persists across the full 2006-2024 period, demonstrating that aggression causally affects performance.

Cross-sectional evidence shows clear temporal erosion: correlations between composite aggression and context-adjusted performance (WAR) declined from $r = 0.197$ ($p = 0.007$) in 2006-2011 to essentially zero ($r = -0.018$, $p = 0.790$) in 2018-2024. Whether the causal effect itself has eroded remains less certain. Within-coach fixed effects show suggestive patterns—point estimates declining 49% from 0.258 games in early to 0.131 games in late eras—but all era-specific effects are statistically non-significant and era-stratified analyses suffer from severe power limitations (27%, 21%, and 13% power respectively), preventing definitive conclusions about temporal changes in the causal effect.

The cross-sectional erosion reflects two concurrent, statistically robust dynamics. First, diffusion: mean aggression increased significantly over time ($\beta = +0.125\%$ per year, $R^2 = 0.83$, $p < 0.001$), rising from -0.35% in the early era to $+1.23\%$ in the late era, a 1.58 percentage point increase ($t = -8.76$, $p < 0.001$, Cohen's $d = 5.37$). As analytics-driven tactics spread league-wide, they ceased to differentiate winners from losers in cross-sectional comparisons. Second, increased variance: rather than converging to a single optimum, coaches diverged—composite aggression standard deviation increased 23% from early to late era. The mechanisms driving erosion remain ambiguous: defensive adaptation, competitive equilibration (aggression becoming table stakes), or other unmeasured factors.

Coaching background matters for stability but not for performance. Both offensive and defensive head coaches show similar weak overall relationships between aggression and performance. Pass-heavy aggression shows statistically significant but small correlations with WAR for both groups ($r = 0.136$, $p = 0.016$ for offensive; $r = 0.146$, $p = 0.021$ for defensive), while composite aggression shows non-significant relationships for both ($r = 0.074$, $p = 0.192$ for offensive; $r = 0.120$, $p = 0.058$ for defensive). However,

background profoundly affects stability: offensive coaches show nearly double the year-to-year persistence in fourth down philosophy compared to defensive coaches ($r = 0.690$ vs $r = 0.373$). For offensive coaches, aggression—especially fourth down decisions—represents deeply held strategic beliefs that persist across contexts. For defensive coaches making offensive decisions, aggression is more contextual and reactive, adapting to personnel and game situations.

The popular narrative of coaching “trees” passing down strategic DNA from mentors to protégés finds minimal empirical support for aggression. Overall composite aggression inheritance is essentially zero ($r = 0.050$, $p = 0.261$). Fourth down philosophy provides the sole exception: it shows modest but statistically significant inheritance ($r = 0.150$, $p < 0.001$), driven entirely by offensive mentors ($r = 0.217$, $p < 0.001$) with defensive mentors transmitting nothing ($r = -0.016$, $p = 0.808$). Among coordinators, offensive coordinators transmit both fourth down philosophy ($r = 0.191$, $p = 0.008$) and pass-heavy tendencies ($r = 0.200$, $p = 0.005$) to protégés who become head coaches. Defensive coordinators transmit no detectable strategic tendencies ($r \approx 0$ for all components).

These findings reveal a fundamental tension in competitive environments: the strategies that create advantage inevitably destroy it through their own success. For NFL organizations, this means hiring aggressive coaches no longer improves outcomes—everyone now executes what was once cutting-edge analytics. The critical skill has shifted from tactical execution to strategic foresight: identifying tomorrow’s inefficiency before it becomes common knowledge. Pedigree-based hiring compounds this problem, as coaching trees transmit yesterday’s innovations rather than generating new ones. Fourth down philosophy represents the lone exception where beliefs persist and transmit, making it the only aggression dimension worth probing in candidate assessment.

The implications for team-building are immense. Organizations investing millions in analytics departments to identify optimal fourth-down thresholds and two-point conversion situations are optimizing strategies that no longer differentiate winners from losers. The competitive frontier has moved. Success now requires coaches who can identify the next market inefficiency—perhaps in player evaluation, in-game adaptation speed, or scheme innovation—before it becomes obvious to all thirty-two teams. In the NFL’s relentless competitive environment, yesterday’s cutting edge is today’s baseline. The advantage lies not in being aggressive, but in being first to the next source of advantage.

Acknowledgment: This analysis builds on play-by-play data from `nfl_data_py` and coaching history data scraped from Pro Football Reference. Performance metrics (WAR) were calculated separately. All code and detailed methodology are available in the project repository.

References

- Alamar, B. C. and Mehrotra, V. (2006). Beyond ‘Moneyball’: The rapidly evolving world of sports analytics. *Analytics Magazine*, 22–24.
- Benjamini, Y. and Hochberg, Y. (1995). Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B (Methodological)*, 57(1), 289–300.
- Burke, B. (2009). Fourth down study. *Advanced NFL Stats* (later renamed *Advanced Football Analytics*).
- Carroll, B., Palmer, J., and Thorn, P. (2009). The hidden game of football. *Total Access*.
- CBS Sports (2021, October 18). Brandon Staley’s fourth-down decisions burn Chargers as head coach tried to generate life into dormant offense. *CBS Sports*.
- Chen, X. and Chau, M. C. (2008). The role of analytics in professional sports: The case of the NFL. *Communications of the ACM*, 51(1), 120–121.
- Goff, B. L. and Wisley, T. L. (2006). Is there a managerial life cycle? Evidence from the NFL. *Managerial and Decision Economics*, 27(6), 563–572.
- Horowitz, I. and Shim, R. (2014). The fourth down puzzle in the National Football League. *Journal of the Royal Statistical Society: Series A*, 177(3), 717–733.
- Los Angeles Chargers (2023, December 15). Chargers Part Ways with Head Coach Brandon Staley and General Manager Tom Telesco. Retrieved from <https://www.chargers.com/news/>

- NBC Sports (2023, September 25). Chargers' Brandon Staley defends aggressive fourth-down call in loss to Vikings. *NBC Sports*.
- Rogers, E. M. (2007). *Diffusion of Innovations* (5th ed.). Free Press, New York.
- Romer, D. (2006). Do firms maximize? Evidence from professional football. *Journal of Political Economy*, 114(2), 340–365.
- Schatz, A. (2003). DVOA (Defense-adjusted Value Over Average) metrics. *Football Outsiders*.
- Sportskeeda (2024, January 31). How much do NFL head coaches make? Exploring 5 highest and 5 lowest-paid coaches in 2024. *Sportskeeda*. Retrieved from <https://www.sportskeeda.com/nfl/>
- Williamson, J. (2026). Coach WAR: A counterfactual approach to evaluating NFL head coach impact in the modern era. *[Journal TBD]*.
- Yam, D. and Lopez, M. J. (2019). What was lost? A causal estimate of fourth down behavior in the National Football League. *Journal of Sports Economics*, 20(4), 479–501.